

Report of Q5

Smoke Status Recognition

50015627 JIANG Zhuoyang

1. Question Description:

Smoking is one of the major health problems. From a biomedical point of view, we can determine whether a patient smokes from certain biometric information. I was required to implement a binary algorithm to predict a patient's smoking status given information about various other health indicators.

- Load Data
- Preprocessing and EDA
- Feature Engineering
- Model Training
- Use Model to Do Classification On Test set

2. Load Data Especially Video Data

Load Training Data:

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 159256 entries, 0 to 159255
Data columns (total 24 columns):
#   Column                Non-Null Count  Dtype
---  -
0   id                    159256 non-null  int64
1   age                  159256 non-null  int64
2   height(cm)          149701 non-null  float64
3   weight(kg)           159256 non-null  int64
4   waist(cm)            159248 non-null  float64
5   eyesight(left)       149692 non-null  float64
6   eyesight(right)      159247 non-null  float64
7   hearing(left)        159256 non-null  int64
8   hearing(right)       149701 non-null  float64
9   systolic             159256 non-null  int64
10  relaxation            159256 non-null  int64
11  fasting blood sugar  159256 non-null  int64
12  Cholesterol           159256 non-null  int64
13  triglyceride          159256 non-null  int64
14  HDL                  159256 non-null  int64
15  LDL                  159256 non-null  int64
16  hemoglobin            159256 non-null  float64
17  Urine protein         149669 non-null  float64
18  serum creatinine     159256 non-null  float64
19  AST                  159256 non-null  int64
20  ALT                  159256 non-null  int64
21  Gtp                  159256 non-null  int64
22  dental caries        159256 non-null  int64
23  smoking              159256 non-null  int64
dtypes: float64(8), int64(16)
memory usage: 29.2 MB
```

Figure.1 Training Data Information

Load Test Data:

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 106171 entries, 0 to 106170
Data columns (total 23 columns):
#   Column                Non-Null Count  Dtype
---  -
0   id                    106171 non-null  int64
1   age                  106171 non-null  int64
2   height(cm)          106171 non-null  int64
3   weight(kg)           106171 non-null  int64
4   waist(cm)            106171 non-null  float64
5   eyesight(left)       106171 non-null  float64
6   eyesight(right)      106171 non-null  float64
7   hearing(left)        106171 non-null  int64
8   hearing(right)       106171 non-null  int64
9   systolic             106171 non-null  int64
10  relaxation            106171 non-null  int64
11  fasting blood sugar  106171 non-null  int64
12  Cholesterol           106171 non-null  int64
13  triglyceride          106171 non-null  int64
14  HDL                  106171 non-null  int64
15  LDL                  106171 non-null  int64
16  hemoglobin            106171 non-null  float64
17  Urine protein         106171 non-null  int64
18  serum creatinine     106171 non-null  float64
19  AST                  106171 non-null  int64
20  ALT                  106171 non-null  int64
21  Gtp                  106171 non-null  int64
22  dental caries        106171 non-null  int64
dtypes: float64(5), int64(18)
memory usage: 13.6 MB
```

Figure.2 Test Data Information

3. Preprocessing and EDA

(1) Missing Processing:

Fill missing values using the mean

```
1 # Fill missing values using the mean
2 imputer = SimpleImputer(strategy='mean')
3 X_imputed = imputer.fit_transform(X)
4 X_test_imputed = imputer.fit_transform(X_test)
```

Figure.3 Fill missing values using the mean

(2) EDA

Draw box plots and a Correlation Heatmap for all features, we can find that there are features in different distributions, and some of them have relationship we can do different kind of Feature Engineering on features with different distributions and use the relationship between features to generate new features.

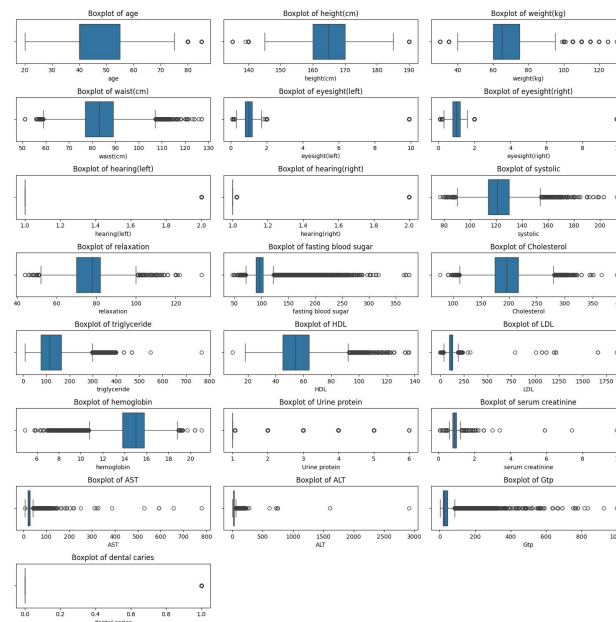


Figure.4 Draw a box plot of all features




Figure.5 Correlation Heatmap of Features

4. Feature Engineering

Create with feature as below:

- Create new feature - BMI:

$$\text{BMI} = \text{weight_kg} / ((\text{height_cm} / 100)^2)$$

- Create new feature - Average Eyesight

$$\text{average_eyesight} = (\text{left_eyesight} + \text{right_eyesight}) / 2$$

- Create new feature - Average Hearing

$$\text{average_hearing} = (\text{left_hearing} + \text{right_hearing}) / 2$$

Do Feature Engineering for both Training and Test data:

```
1 # Concatenate training and test data without labels
2 all_data = pd.concat([train_data.iloc[:, 1:-1], test_data.iloc[:, 1:-1]], axis=0)
3
4 # Impute missing values using mean strategy
5 imputer = SimpleImputer(strategy='mean')
6 all_data_imputed = imputer.fit_transform(all_data)
7
8 # Create new feature: BMI
9 height_cm = all_data_imputed[:, all_data.columns.get_loc('height(cm)')]
10 weight_kg = all_data_imputed[:, all_data.columns.get_loc('weight(kg)')]
11 BMI = weight_kg / ((height_cm / 100) ** 2) # BMI calculation formula
12 all_data_imputed = np.column_stack((all_data_imputed, BMI.reshape(-1, 1)))
13
14 # Create new feature: Average Eyesight
15 left_eyesight = all_data_imputed[:, all_data.columns.get_loc('eyesight(left)')]
16 right_eyesight = all_data_imputed[:, all_data.columns.get_loc('eyesight(right)')]
17 average_eyesight = (left_eyesight + right_eyesight) / 2 # Average eyesight calculation formula
18 all_data_imputed = np.column_stack((all_data_imputed, average_eyesight.reshape(-1, 1)))
19
20 # Create new feature: Average Hearing
21 left_hearing = all_data_imputed[:, all_data.columns.get_loc('hearing(left)')]
22 right_hearing = all_data_imputed[:, all_data.columns.get_loc('hearing(right)')]
23 average_hearing = (left_hearing + right_hearing) / 2 # Average hearing calculation formula
24 all_data_imputed = np.column_stack((all_data_imputed, average_hearing.reshape(-1, 1)))
```

Figure.7 Input Organization of 2D Convolution Deep Model

The Number of features after feature engineering shows as below:

Number of features in training set: 25

Number of features in testing set: 25

Figure.7 Number of features after feature engineering

5. Model Training and Evaluation

(1) Method Description - LightGBM :

The LGBMClassifier is a classification algorithm provided by the LightGBM library, which stands for Light Gradient Boosting Machine. It's based on the gradient boosting framework and is specifically designed for handling large datasets efficiently while delivering high performance.

Specifically, LightGBM's Gradient Boosting Framework means it constructed an ensemble of decision trees sequentially to minimize the errors made by the preceding trees.

It's optimized for speed and memory efficiency using a histogram-based approach for binning continuous feature values and making split decisions vertically on these bins, reducing memory usage and speeding up training.

Unlike traditional depth-wise tree growth, LightGBM grows trees leaf-wise. It selects the leaf with the maximum delta loss to grow, allowing for a more complex tree and better capture of feature interactions.

What's more LightGBM can handle categorical features directly, avoiding the need for one-hot encoding and reducing memory overhead.

(2) Model Construction and GridSearch Defining:

Define the LGBMClassifier model and apply a parameters grid for automatically tuning, the GridSearch Method will help us find a set of good hyperparameter.

```
1 # Define the parameters grid for tuning
2 param_grid = {
3     'num_leaves': [20, 30, 40],
4     'learning_rate': [0.05, 0.1, 0.2],
5     'max_depth': [5, 10, -1], # -1 means no limit
6     'min_child_samples': [20, 30, 50],
7 }
8
9 # Split the dataset into training and validation sets
10 X_train_improved, X_val_improved, y_train, y_val = train_test_split(X_improved, y, test_size=0.01, random_state=42)
11
12 # Initialize LightGBM classifier
13 lgb_classifier = lgb.LGBMClassifier(boosting_type='gbdt', objective='binary', metric='auc', n_jobs=-1, verbosity=1)
14
15 # Perform GridSearchCV
16 grid_search = GridSearchCV(estimator=lgb_classifier, param_grid=param_grid, scoring='roc_auc', cv=5)
17
```

Figure.8 Model Construction and GridSearch

(3) Training Processing:

The training processing of LGBMClassifier can be should with the information:

```
In [25]: 1 # Fit GridSearchCV to find best parameters
2         grid_search.fit(X_train_improved, y_train)
3
4         # Get the best model from GridSearchCV
5         best_params = grid_search.best_params_
6         best_score = grid_search.best_score_
7         best_model = grid_search.best_estimator_
8
9         # Print the best parameters and score
10        print('Best parameters: ', best_params)
11        print('Best score: ', best_score)
12
13        [LightGBM] [Info] Total Bins 2357
14        [LightGBM] [Info] Number of data points in the train set: 126131, number of used features: 25
15        [LightGBM] [Info] [binary:BoostFromScore]: pavg=0.437117 -> initscore=-0.252871
16        [LightGBM] [Info] Start training from score -0.252871
17        [LightGBM] [Info] Number of positive: 55134, number of negative: 70997
18        [LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.008740 seconds.
19        You can set 'force_col_wise=true' to remove the overhead.
20        [LightGBM] [Info] Total Bins 2364
21        [LightGBM] [Info] Number of data points in the train set: 126131, number of used features: 25
22        [LightGBM] [Info] [binary:BoostFromScore]: pavg=0.437117 -> initscore=-0.252871
23        [LightGBM] [Info] Start training from score -0.252871
24        [LightGBM] [Info] Number of positive: 68917, number of negative: 88746
25        [LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.011090 seconds.
26        You can set 'force_col_wise=true' to remove the overhead.
27        [LightGBM] [Info] Total Bins 2397
28        [LightGBM] [Info] Number of data points in the train set: 157663, number of used features: 25
29        [LightGBM] [Info] [binary:BoostFromScore]: pavg=0.437116 -> initscore=-0.252875
30        [LightGBM] [Info] Start training from score -0.252875
31        [LightGBM] [Warning] No further splits with positive gain, best gain: -inf
```

Figure.9 Model Construction and GridSearch

(4) Tuning Result and Evaluation:

Best Parameters: {'learning_rate': 0.2, 'max_depth': 10, 'min_child_samples': 50, 'num_leaves': 30}
Best ROC-AUC Score: 0.865002584046359
Validation ROC-AUC Score: 0.8569691415960631

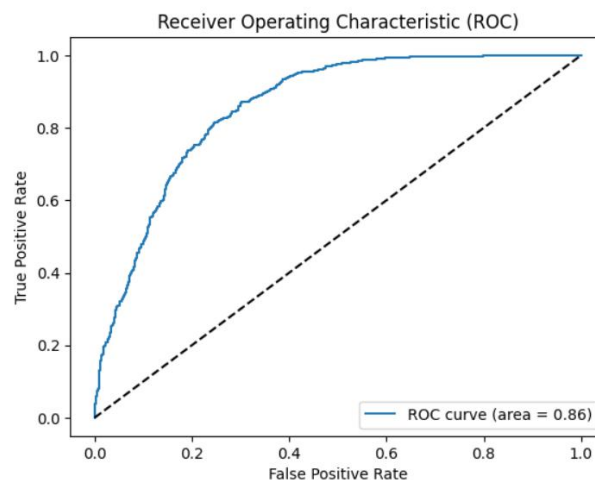


Figure.12 Tuning Result and Evaluation

6. Use Model to Do Classification On Test set

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 106171 entries, 0 to 106170
Data columns (total 2 columns):
#   Column      Non-Null Count  Dtype
---  -
0   id          106171 non-null  int64
1   smoking     106171 non-null  float64
dtypes: float64(1), int64(1)
memory usage: 1.6 MB
```

Figure.13 Output.csv